

Outline: Ultra-High Sample Size Supervised Problems (F. Chiaromonte)

Statistical Methods for Large, Complex Data

- Main reference: R. Zhu, P. Ma, M. Mahoney, B. Yu (2015) Optimal Subsampling Approaches for Large Sample Linear Regression. arXiv 1509.05111v1.
- Additional references: throughout the slides.

A SHORT REMINDER ON THE NON-PARAMETRIC BOOTSTRAP (resampling with replacement)

The bootstrap algorithm:

- Generate B bootstrap samples of size n drawing with replacement from the data

$$x_{(b)}^* \quad b = 1 \dots B$$

- On each, compute the statistic of interest – producing B bootstrap values; “copies” that mimic its sampling variability

$$\hat{\theta}_{(b)}^* = g(x_{(b)}^*) \quad b = 1 \dots B$$

- Estimate the **standard error** as the standard deviation of the bootstrap values


$$\widehat{se}_{BT} = \sqrt{\frac{1}{B-1} \sum_{b=1}^B (\hat{\theta}_{(b)}^* - \hat{\theta}_{(\cdot)}^*)^2} \quad \hat{\theta}_{(\cdot)}^* = \frac{1}{B} \sum_{b=1}^B \hat{\theta}_{(b)}^*$$

Already rather good with **B=200**.

- Build the corresponding empirical cdf \hat{G} , a simulated version of the **sampling distribution**
- Use its percentiles to define a percentile-based $(1-\alpha)$ coverage **confidence interval** as

$$LOW = G^{-1}\left(\frac{\alpha}{2}\right) \quad UP = G^{-1}\left(1 - \frac{\alpha}{2}\right)$$

Not the ultimate answer (can be further improved bias and potential non-constant variance corrections), but good with **B=1000, 2000**.



Rationale 1: since we cannot take more samples of size n from F , we do the next best thing, i.e., generate them from the empirical distribution \hat{F}_n . This is our best estimate of F based on the available sample of n observations (in fact, its non-parametric MLE).

Rationale 2: the bootstrap lets us gauge the variability of $\hat{\theta}$ creating perturbations of the original data by resampling (with replacement). These are less local of the perturbations created with the Jackknife (create n samples of size $n-1$ deleting one observation at a time).

BACK TO ULTRA-HIGH SAMPLE SIZE

Stochastic mechanism:

Y, ε r.vbls in R^1 X r.vct in R^p

$$E(X) = 0 \quad \text{Cov}(X) = \Sigma_X$$

$$\varepsilon \text{ indep } X \quad \varepsilon \sim N(0, \sigma^2)$$

Estimation (fitting)

$$\hat{\beta}_{LS} = (X'X)^{-1}X'Y$$

Independent sampling:

$$Y_{nx1} = X_{n \times p} \beta_{p \times 1} + \varepsilon_{nx1}$$

When n is very large, computationally prohibitive even for relatively small p ; $O(np^2)$

An intuitive solution: **SUBSAMPLING**

(A) Subsample with replacement from the data (empirical distribution; weight of each data point $w_{i,ID}=1/n$)

$$Y^*_{rx1}; X^*_{rxp} \quad r \ll n$$

UNIF

UNIF refers to the fact that in subsampling each original data point has the same weight.

(B) Use the subsample as a surrogate and compute

$$\hat{\beta}^*_{LS}(w_{ID}) = (X^{*'} X^*)^{-1} X^{*'} Y^*$$

- Approximating the full LS vector:
(Conditionally on the data) $E(\hat{\beta}^*_{LS}(w_{ID})|F_n) = \hat{\beta}_{LS}$
- Estimating the true effects:
(Unconditionally) $E(\hat{\beta}^*_{LS}(w_{ID})) = \beta$

Problem: **accuracy of the estimator can be pretty bad for moderate r!**

To improve accuracy, we need to **weigh the data points non-uniformly** when subsampling.

How? People have used **leverage scores** – exact or approximate:

$H = X(X'X)^{-1}X'$ Orthogonal projection operator of the LS (Hat-matrix)

$$h_{ii} = x_i(X'X)^{-1}x_i = ||x_i||_{(X'X)^{-1}}^2$$

$$w_i = h_{ii}/p$$

BLEV

Basic Leverage Scores; length (squared) of x_i in the S_X^{-1} inner product.

Rationale: **represent the periphery of the data** when subsampling.

Estimating the leverage scores on the full data is prohibitive (as much as computing the full LS!) so people proposed approximate calculations

ALEV

Some people also “shrink” the leverage weighting to the uniform weighting

SLEV

Non-uniform weights induce a bias in approximating the full LS, which cannot be controlled with r (does not vanish as the subsample size grows)... but it reduces variance in approximating the full LS. In all, accuracy improves.

Note: *they do NOT induce bias in estimating the true effects...* more later on variance/accuracy.

Important results in the paper:

To correct the bias in approximating the full LS, one needs to ***weigh the LS*** on the subsample according to the weights used in subsampling!

$$\hat{\beta}_{WLS}^*(w) = (X^{*'} W X^*)^{-1} X^{*'} W^{-1/2} Y^* \quad W = \text{Diag}(w)$$

Under regularity conditions (bounding high order moments of the X's on the full data with weights), for large enough r

- **Thm 1.** Approximating the full LS vector: (Conditionally on the data) the weighted estimator is approximately Gaussian around the full LS with covariance matrix $V(w)$

$$V^{-1/2}(w)(\hat{\beta}_{WLS}^*(w) - \hat{\beta}_{LS}) \xrightarrow[r \rightarrow \infty]{D} N(0, I)$$

- **Thm 2.** Estimating the true effects: (Unconditionally) the weighted estimator is approximately Gaussian around the true β with covariance matrix $V_o(w) > V(w)$

$$V_o^{-1/2}(w)(\hat{\beta}_{WLS}^*(w) - \beta) \xrightarrow[r \rightarrow \infty]{D} N(0, I)$$

- **Thm 3.** Provided r grows linearly with n (e.g. a given fraction), weighing the LS actually does NOT improve accuracy for estimating the true effects

$$V_o(\hat{\beta}_{LS}^*(w)) \leq V_o(\hat{\beta}_{WLS}^*(w))$$

... FOR ESTIMATING THE TRUE β IT IS BETTER TO WEIGH THE SUBSAMPLING, BUT NOT TO WEIGH THE LS!

Also, the authors demonstrate that **(B,A,S) LEV weighting is NOT optimal in increasing accuracy** for approximating the full LS. The **Optimal weights** have a different form

$$w_i = \frac{\sqrt{(1 - h_{ii})} ||x_i||}{\sum \sqrt{(1 - h_{jj})} ||x_j||} = \frac{\sqrt{(1 - ||x_i||^2_{(X'X)^{-1}}) ||x_i||^2}}{\sum \sqrt{(1 - ||x_j||^2_{(X'X)^{-1}}) ||x_j||^2}} \quad \boxed{\text{OPT}}$$

Assuming the data is not too dispersed (compact bulk, not important how we account for the periphery) so that the leverages are not too heterogeneous (all $\sim p/n$) we can simply use what they call **Predictor Length weights**

$$w_i = \frac{||x_i||}{\sum ||x_j||} \quad \boxed{\text{PL}}$$

Similar to the LEV weights – but using the standard inner product!
Much cheaper to compute since they do not involve S^{-1} ; $O(np)$.

... FOR APPROXIMATING THE FULL LS, SUBSAMPLE WITH PL WEIGHTS AND USE THE PL-WEIGHTED ESTIMATOR.

... FOR ESTIMATING THE TRUE β , SUBSAMPLE WITH PL WEIGHTS AND USE THE UNWEIGHTED ESTIMATOR.

SIMULATIONS TO EVALUATE PERFORMANCE ON “FINITE SUBSAMPLES”

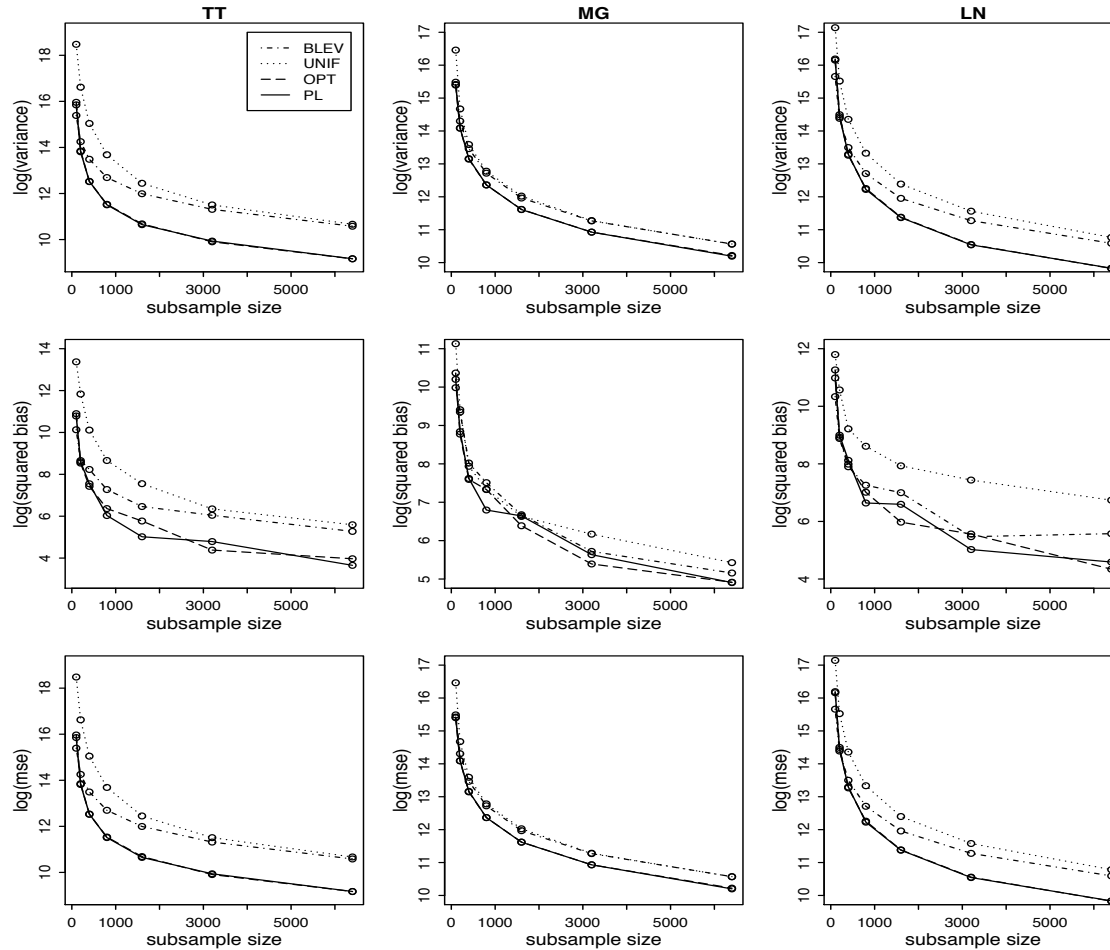
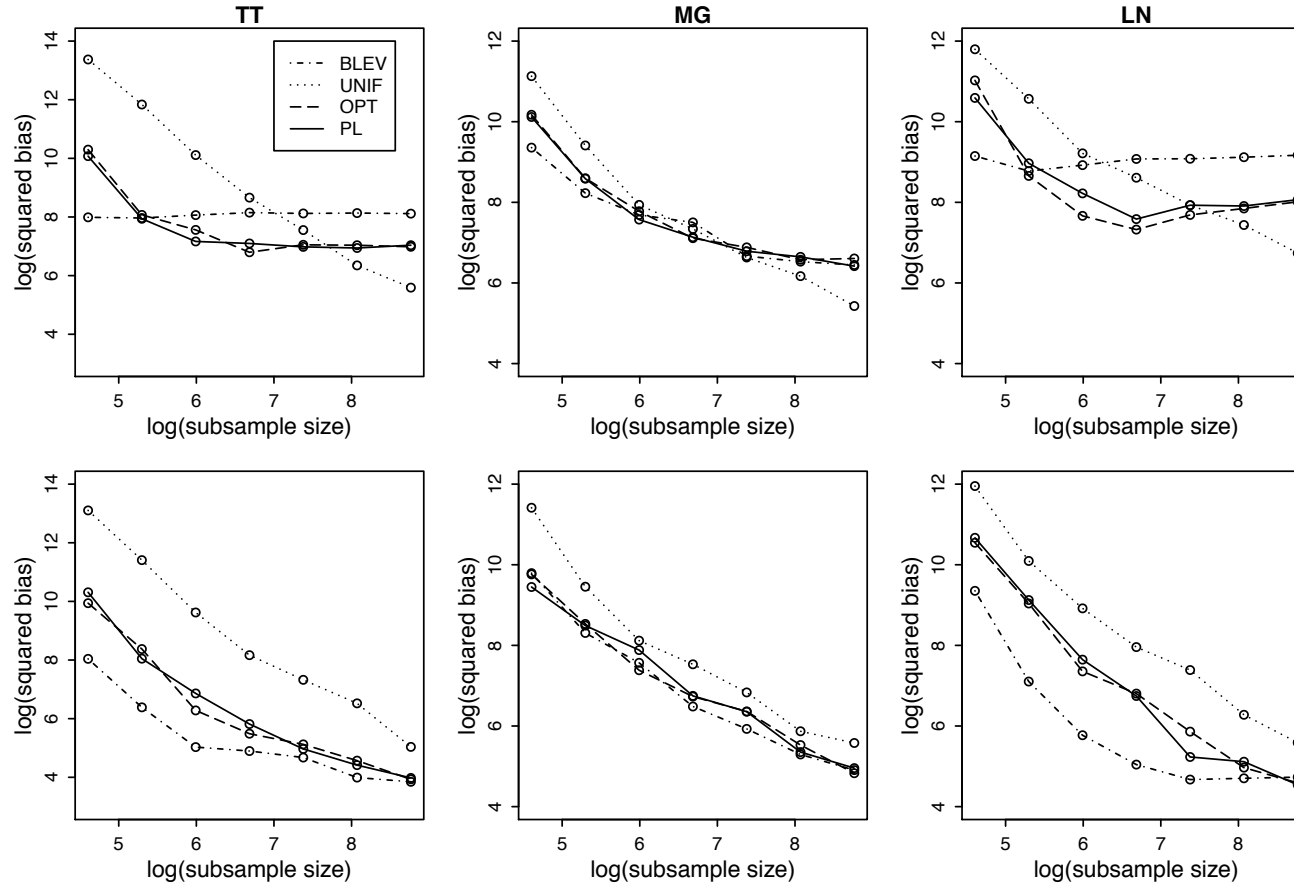


Figure 4: Empirical variances, squared biases and mean-squared error of $\mathbf{X}\tilde{\beta}$ for predicting $\mathbf{X}\hat{\beta}_{ols}$. Upper panels are the logarithm of variances, middle panels the logarithm of squared bias, and lower panels the logarithm of mean-squared error. From left to right: TT, MG and LN data, respectively.

NOTE: here empirical variances, squared biases and MSEs from simulations are on response predictions, not β estimates.



SIMULATIONS TO EVALUATE
PERFORMANCE ON “FINITE SUBSAMPLES”

NOTE: here empirical variances, squared
biases and MSEs from simulations are on
response predictions, not β estimates.

Figure 5: Empirical squared biases of $\mathbf{X}\tilde{\beta}^u$ for predicting $\mathbf{X}\hat{\beta}_{ols}$ and $\mathbf{X}\beta$, respectively. From top to bottom: upper panels are results for predicting $\mathbf{X}\hat{\beta}_{ols}$, and lower panels results for predicting $\mathbf{X}\beta$. From left to right: TT, MG and LN data, respectively.

ADDITIONAL REFERENCE:

A similar analysis is provided for **GLM**, logistic regression with ultra-high sample sizes:

H. Wang, R. Zhu, P. Ma (2018) Optimal Subsampling for Large Sample Logistic Regression. JASA, 113(522): 829–844.

Abstract: For massive data, the family of subsampling algorithms is popular to downsize the data volume and reduce computational burden. Existing studies focus on approximating the ordinary least squares estimate in linear regression, where statistical leverage scores are often used to define subsampling probabilities. In this paper, we propose fast subsampling algorithms to efficiently approximate the maximum likelihood estimate in logistic regression. We first establish consistency and asymptotic normality of the estimator from a general subsampling algorithm, and then derive optimal subsampling probabilities that minimize the asymptotic mean squared error of the resultant estimator. An alternative minimization criterion is also proposed to further reduce the computational cost. The optimal subsampling probabilities depend on the full data estimate, so we develop a two-step algorithm to approximate the optimal subsampling procedure. This algorithm is computationally efficient and has a significant reduction in computing time compared to the full data approach. Consistency and asymptotic normality of the estimator from a two-step algorithm are also established. Synthetic and real data sets are used to evaluate the practical performance of the proposed method.

An alternative to subsampling: **SPLIT-AND-CONQUER, PARTITIONING STRATEGIES.**

Note: sometimes the data cannot be (physically) merged in one “location” to perform a computation.

For a given procedure one needs to

- split
- perform peripheral computation
- and then re-merge in a statistically appropriate way.

Some references:

- J. Fan, F. Han, H. Liu (2014) Challenges of big data analysis. National Science Review.
- A. Kleiner, A. Talwalkar, P. Sarkar, M. Jordan (2014) A scalable bootstrap for massive data. JRSS-B.
- X. Chen, M. Xie (2014) A split-and-conquer approach for analysis of extraordinarily large data. Statistica Sinica.
- Y. Zhang, J. Duchi, M. Weinwright (2015) Divide and conquer kernel Ridge regression: a distributed algorithm with minimax optimal rates. Journal of Machine Learning Research.

adding large p , perhaps in some of the parts...

- H. Battey, J. Fan, H. Liu, J. Lu, Z. Zhu (2018) Distributed testing and estimation under sparse high dimensional models. AoS.
- L. Tang, L. Zho, P. Song (2020) Distributed Simultaneous Inference in Generalized Linear Models via Confidence Distribution. Journal of Multivariate Analysis.

COMPUTER SCIENCE

Challenges of Big Data analysis

Jianqing Fan^{1,*}, Fang Han² and Han Liu¹

ABSTRACT

Big Data bring new opportunities to modern society and challenges to data scientists. On the one hand, Big Data hold great promises for discovering subtle population patterns and heterogeneities that are not possible with small-scale data. On the other hand, the massive sample size and high dimensionality of Big Data introduce unique computational and statistical challenges, including scalability and storage bottleneck, noise accumulation, spurious correlation, incidental endogeneity and measurement errors. These challenges are distinguished and require new computational and statistical paradigm. This paper gives overviews on the salient features of Big Data and how these features impact on paradigm change on statistical and computational methods as well as computing architectures. We also provide various new perspectives on the Big Data analysis and computation. In particular, we emphasize on the viability of the sparsest solution in high-confidence set and point out that exogenous assumptions in most statistical methods for Big Data cannot be validated due to incidental endogeneity. They can lead to wrong statistical inferences and consequently wrong scientific conclusions.

Keywords: Big Data, noise accumulation, spurious correlation, incidental endogeneity, data storage, scalability

A SPLIT-AND-CONQUER APPROACH FOR ANALYSIS OF EXTRAORDINARILY LARGE DATA

Xueying Chen and Min-ge Xie

Rutgers University

Abstract: If there are datasets, too large to fit into a single computer or too expensive for a computationally intensive data analysis, what should we do? We propose a *split-and-conquer* approach and illustrate it using several computationally intensive penalized regression methods, along with a theoretical support. We show that the split-and-conquer approach can substantially reduce computing time and computer memory requirements. The proposed methodology is illustrated numerically using both simulation and data examples.

Key words and phrases: Big data, combining results from independent analyses, distributed computing, generalized linear models, large sample theory, penalized regression.

Distributed Simultaneous Inference in Generalized Linear Models via Confidence Distribution

Lu Tang^a, Ling Zhou^b, Peter X.-K. Song^{c,*}

^a*Department of Biostatistics, University of Pittsburgh, Pittsburgh, PA 15261, USA*

^b*Center of Statistical Research, Southwestern University of Finance and Economics, Chengdu, Sichuan, China*

^c*Department of Biostatistics, University of Michigan, Ann Arbor, MI 48109, USA*

Abstract

We propose a distributed method for simultaneous inference for datasets with sample size much larger than the number of covariates, i.e., $N \gg p$, in the generalized linear models framework. When such datasets are too big to be analyzed entirely by a single centralized computer, or when datasets are already stored in distributed database systems, the strategy of divide-and-combine has been the method of choice for scalability. Due to partition, the sub-dataset sample sizes may be uneven and some possibly close to p , which calls for regularization techniques to improve numerical stability. However, there is a lack of clear theoretical justification and practical guidelines to combine results obtained from separate regularized estimators, especially when the final objective is simultaneous inference for a group of regression parameters. In this paper, we develop a strategy to combine bias-corrected lasso-type estimates by using confidence distributions. We show that the resulting combined estimator achieves the same estimation efficiency as that of the maximum likelihood estimator using the centralized data. As demonstrated by simulated and real data examples, our divide-and-combine method yields nearly identical inference as the centralized benchmark.

Keywords: Bias correction, Confidence distribution, Inference, Lasso, Meta-analysis, Parallel computing.

2010 MSC: Primary 62H15, Secondary 62F12
